

# A Provenance-aware memory object model for C

## ISO/IEC TS 6010:2023

Jens Gustedt

INRIA – Camus

ICube – ICPS

Université de Strasbourg joint work with

Peter Sewell, Kayvan Memarian, Victor B. F. Gomes, Martin Uecker

ISO/IEC JTC 1/SC 22/WG14 N3057

<https://open-std.org/JTC1/SC22/WG14/www/docs/3057.pdf>



<https://modernc.gforge.inria.fr/>



# Table of Contents

- 1 Motivation
- 2 Provenance based aliasing analysis
- 3 Standardization



# The C programming language

## ISO/IEC IS 9899:2018

### Features

**good:** standardized, portable, stable, simple, efficient, extensible

**bad:** safety, out-of-bounds access, lack of encapsulation, pointer casts, lack of composability, side channels

### One of the major programming languages

C is the base and interface for

- all modern processor architectures
  - 16 bit microprocessor (10 ¢)
  - ...
  - 128 bit multi-core high performance processor
- many other programming languages

# C is highly optimizable

## "As-if" rule

- The compiled code only performs visible effects.
- How the compiler achieves that is their business/secret.
- Internal rewriting into more efficient code with equivalent semantics.

## Storage-backed objects

- Pointers allow access to objects from different angles/views.
- Many optimizations can only be applied if access is unique.
- Possible aliasing inhibits optimizations.



# Aliasing techniques in C

## Type-based aliasing analysis

```
void f(int* a, float* b) {  
    // *a and *b will never alias  
}
```

## restrict-based aliasing analysis

```
void g(float*restrict a, float*restrict b) {  
    // *a and *b will never alias  
}
```

## interdiction based aliasing prevention

```
// address of toto may never be taken  
register double toto = 35.4;
```

# How to avoid aliasing for out-of-bounds accesses?

## Basic example

```
signed int y = 2, x = 1;
int main() {
    signed int* p = &x + 1; // one beyond storage
    signed int* q = &y;      // pointers ok
    if (!memcmp(&p, &q, sizeof(p))) { // exposure!
        *p = 11; // formed from x, accessing y?
    }
}
```

## Question

- How to complement C's memory model such that the access to `*p` is forbidden?
- How to build tools that analyze such situations?

# Table of Contents

- 1 Motivation
- 2 Provenance based aliasing analysis
- 3 Standardization



# Definitions

## Storage instance

A storage instance is a continuous **chunk of storage** with a **unique ID**

- When a declared object or compound literal is instantiated
  - Static storage duration:                      program startup → program termination
  - Thread storage duration:                      thread startup → thread termination
  - Automatic storage duration:                      scope entrance → scope exit
- Allocated storage duration:
  - { **malloc**, **calloc**, **realloc** } → { **free**, **realloc** }
- Recycled storage (stack or heap) → **new ID**

## Provenance


Each valid pointer holds

- an ID of the corresponding storage instance
- an address



# Detection of conflicts

## A view from Cerberus



<http://cerberus.cl.cam.ac.uk/cerberus>

# Exposure and synthesis

## Conversion between pointers and integers

Bad, bad, bad!

All type information is lost!

All provenance information is lost!

## Track provenance?

- through integers?
- through representation bytes?
- through control flow?



# Exposure and synthesis

## Terminology

- pointer to integer conversion: the whole storage instance is **exposed**
- integer to pointer conversion: the pointer is **synthesized**

## Similar features

- **printf** or **scanf** with %p
- byte fiddling with in-storage representation of pointers
- **fwrite** or **fread** of a pointer representation



# Exposure and synthesis

## Rule

$J$  has been exposed  $\rightarrow$  A pointer to  $J$  may be synthesized

## Interdiction

If a pointer to  $J$  is synthesized and  $J$  has not been exposed:


undefined behavior



# Disambiguation

## Two provenances for a synthesized pointer

A synthesized pointer **p** may point simultaneously to the end of exposed **J** and the start of exposed **K**



## Rule

Only a use with either provenance **J** or **K** is allowed.



# Table of Contents

- 1 Motivation
- 2 Provenance based aliasing analysis
- 3 Standardization



# The full technical specification

## ISO/IEC TS 6010:2023

<https://open-std.org/JTC1/SC22/WG14/www/docs/3057.pdf>

# A first step: know your constituency

What do programmers think?

What do compiler implementors think?

What do experts in WG14 think?

All depends!

- there are more diverging opinions than people
- opinions can be shifted with good arguments and a lot of patience
- everything flows, but very, very slowly

Tedious community work by the Cambridge group around Peter Sewell



# A second step: condense different strategies into sound models

## Questions?

- Which granularity for the model? (allocation, object, member?)
- How to deal with reallocations?
- How to deal with type changes?
- How to deal with information leakage?

## Method

- Yearlong discussions in WG14
- Liaison activities with WG21 (C++)
- A new Memory Model Study Group (head Peter Sewell)
- Mathematical specifications of several models
- An online test tool: Cerberus

# A third step: integration and acceptance

## The C standard itself

- Some new concepts and definitions
- Detailed changes all over
- More than 100 pages impacted

## The community

- Convince implementors to obey to the "new" rules
- Tag behavior that is not conforming in open implementations as bugs
- Have these bugs accepted by the community

## WG14

- Convince them of the proposed changes
- Partial success → propose a technical specification (TS 6010)

# A fourth step: write an ISO norm

## Collect all the material

(weird copyright rules)

- A paper on executable examples
- A paper on the semantic specification
- The "diff" to the C standard

→ three annexes to the TS

## Write the specification itself

- ISO language is quite particular
  - can, may, shall, should . . . all have very restricted use
  - terminology sometimes fixed in other ISO documents
- ISO has no modern collaborative tools
- ISO is quite restrictive for the editorial "style"

# A fifth step: publish an ISO norm

## Navigate it through the voting process

- WG14
- National Bodies (AFNOR, DIN, INCITS, ...)

## publish

- ISO